

SIGNAL MASK

- A process initially inherits the parent's signal mask when it is created, but any pending signals for the parent process are not passed on. A process may query or set its signal mask via the sigprocmask API:

```
#include <signal.h>
```

```
int sigprocmask(int cmd, const sigset_t *new_mask, sigset_t  
*old_mask);
```

Returns: 0 if OK, 1 on error

- The `new_mask` argument defines a set of signals to be set or reset in a calling process signal mask, and the `cmd` argument specifies how the `new_mask` value is to be used by the API.
- The possible values of `cmd` and the corresponding use of the `new_mask` value are:

Cmd value	Meaning
<code>SIG_SETMASK</code>	Overrides the calling process signal mask with the value specified in the <code>new_mask</code> argument.
<code>SIG_BLOCK</code>	Adds the signals specified in the <code>new_mask</code> argument to the calling process signal mask.
<code>SIG_UNBLOCK</code>	Removes the signals specified in the <code>new_mask</code> argument from the calling process signal mask.

- If the actual argument to `new_mask` argument is a NULL pointer, the `cmd` argument will be ignored, and the current process signal mask will not be altered.
- If the actual argument to `old_mask` is a NULL pointer, no previous signal mask will be returned.
- The `sigset_t` contains a collection of bit flags.

The BSD UNIX and POSIX.1 define a set of API known as sigsetops functions which set, reset and query the presence of signals in a sigset_t typed variable:

```
#include<signal.h>

int sigemptyset (sigset_t* sigmask);
int sigaddset (sigset_t* sigmask, const int sig_num);
int sigdelset (sigset_t* sigmask, const int sig_num);
int sigfillset (sigset_t* sigmask);
int sigismember (const sigset_t* sigmask, const int sig_num);
```

- The `sigemptyset` API clears all signal flags in the `sigmask` argument.
 - The `sigaddset` API sets the flag corresponding to the `signal_num` signal in the `sigmask` argument.
 - The `sigdelset` API clears the flag corresponding to the `signal_num` signal in the `sigmask` argument.
 - The `sigfillset` API sets all the signal flags in the `sigmask` argument.
- [all the above functions return 0 if OK, -1 on error]
- The `sigismember` API returns 1 if flag is set, 0 if not set and -1 if the call fails.

The following example checks whether the SIGINT signal is present in a process signal mask and adds it to the mask if it is not there.

```
#include<stdio.h>
#include<signal.h>
int main()
{
    sigset_t    sigmask;
    sigemptyset(&sigmask);          /*initialise set*/

    if(sigprocmask(0,0,&sigmask)==-1)    /*get current signal mask*/
    {
        perror("sigprocmask");
        exit(1);
    }
    else sigaddset(&sigmask,SIGINT); /*set SIGINT flag*/

    sigdelset(&sigmask, SIGSEGV);      /*clear SIGSEGV flag*/
    if(sigprocmask(SIG_SETMASK,&sigmask,0)==-1)
        perror("sigprocmask");
}
```

- A process can query which signals are pending for it via the sigpending API:

```
#include<signal.h>
```

```
int sigpending(sigset_t* sigmask);
```

Returns 0 if OK, -1 if fails.

- The sigpending API can be useful to find out whether one or more signals are pending for a process and to set up special signal handling methods for these signals before the process calls the sigprocmask API to unblock them.
- The following example reports to the console whether the SIGTERM signal is pending for the process:

```
#include<iostream.h>
#include<stdio.h>
#include<signal.h>
int main()
{
    sigset_t sigmask;
    sigemptyset(&sigmask);
    if(sigpending(&sigmask)==-1)
        perror("sigpending");
    else cout << "SIGTERM signal is:" << (sigismember(&sigmask,SIGTERM) ?
    "Set" : "No Set") << endl;
}
```


UNIX also supports following APIs for signal mask manipulation:

```
#include<signal.h>
```

```
int sighold(int signal_num);
```

```
int sigrelse(int signal_num);
```

```
int sigignore(int signal_num);
```

```
int sigpause(int signal_num);
```

SIGACTION

- The sigaction API blocks the signal it is catching allowing a process to specify additional signals to be blocked when the API is handling a signal.

The sigaction API prototype is:

```
#include<signal.h>
```

```
int sigaction(int signal_num, struct sigaction* action, struct  
sigaction* old_action);
```

Returns: 0 if OK, 1 on error

- The struct sigaction data type is defined in the <signal.h> header as:

struct sigaction

{

void (*sa_handler)(int);

sigset_t sa_mask;

int sa_flag;

}

The following program illustrates the uses of sigaction:

```
#include<iostream.h>
```

```
#include<stdio.h>
```

```
#include<unistd.h>
```

```
#include<signal.h>
```

```
void callme(int sig_num)
```

```
{
```

```
    cout<<"catch signal:"<<sig_num<<endl;
```

```
}
```

```
int main(int argc, char* argv[])
```

```
{
```

```
    sigset_t sigmask;
```

```
    struct sigaction action,old_action;
```

```
    sigemptyset(&sigmask);
```

```
    if(sigaddset(&sigmask,SIGTERM)==-1 || sigprocmask(SIG_SETMASK,&sigmask,0)==-1)
```

```
        perror("set signal mask");
```

```
    sigemptyset(&action.sa_mask);
```

```
    sigaddset(&action.sa_mask,SIGSEGV);
```

```
    action.sa_handler=callme;
```

```
    action.sa_flags=0;
```

```
    if(sigaction(SIGINT,&action,&old_action)==-1)
```

```
        perror("sigaction");
```

```
    pause();
```

```
    cout<<argv[0]<<"exists\n";
```

```
    return 0;
```

```
}
```

THE SIGCHLD SIGNAL AND THE waitpid API

- When a child process terminates or stops, the kernel will generate a SIGCHLD signal to its parent process. Depending on how the parent sets up the handling of the SIGCHLD signal, different events may occur:
 - Parent accepts the **default action of the SIGCHLD signal**:
 - o SIGCHLD does not terminate the parent process.
 - o Parent process will be awakened.
 - o API will return the child's exit status and process ID to the parent.
 - o Kernel will clear up the Process Table slot allocated for the child process.
 - o Parent process can call the waitpid API repeatedly to wait for each child it created.

Parent ignores the SIGCHLD signal:

- o SIGCHLD signal will be discarded.
- o Parent will not be disturbed even if it is executing the waitpid system call.
- o If the parent calls the waitpid API, the API will suspend the parent until all its child processes have terminated.
- o Child process table slots will be cleared up by the kernel.
- o API will return a -1 value to the parent process.

Process catches the SIGCHLD signal:

- o The signal handler function will be called in the parent process whenever a child process terminates.
- o If the SIGCHLD arrives while the parent process is executing the waitpid system call, the waitpid API may be restarted to collect the child exit status and clear its process table slots.
- o Depending on parent setup, the API may be aborted and child process table slot not freed.

THE sigsetjmp AND siglongjmp

APIs

- The function prototypes of the APIs are:

```
#include <setjmp.h>
```

```
int sigsetjmp(sigjmp_buf env, int savemask);
```

```
int siglongjmp(sigjmp_buf env, int val);
```


- The `sigsetjmp` and `siglongjmp` are created to support signal mask processing. Specifically, it is implementation dependent on whether a process signal mask is saved and restored when it invokes the `setjmp` and `longjmp` APIs respectively.
- The only difference between these functions and the `setjmp` and `longjmp` functions is that `sigsetjmp` has an additional argument. If `savemask` is nonzero, then `sigsetjmp` also saves the current signal mask of the process in `env`.
- When `siglongjmp` is called, if the `env` argument was saved by a call to `sigsetjmp` with a nonzero `savemask`, then `siglongjmp` restores the saved signal mask. The `siglongjmp` API is usually called from user-defined signal handling functions. This is because a process signal mask is modified when a signal handler is called, and `siglongjmp` should be called to ensure the process signal mask is restored properly when “jumping out” from a signal handling function.

- The following program illustrates the uses of sigsetjmp and siglongjmp APIs.

```
#include<iostream.h>
#include<stdio.h>
#include<unistd.h>
#include<signal.h>
#include<setjmp.h>

sigjmp_buf env;
void callme(int sig_num)
{
    cout<< "catch signal:" <<sig_num <<endl;
    siglongjmp(env,2);
}
int main()
{
    sigset_t sigmask;
    struct sigaction action,old_action;

    sigemptyset(&sigmask);
```

```
if(sigaddset(&sigmask,SIGTERM)==-1) || sigprocmask(SIG_SETMASK,&sigmask,0)==-1)

    perror("set signal mask");

sigemptyset(&action.sa_mask);

sigaddset(&action.sa_mask,SIGSEGV);

action.sa_handler=(void(*)())callme;

action.sa_flags=0;

if(sigaction(SIGINT,&action,&old_action)==-1)

    perror("sigaction");

if(sigsetjmp(env,1)!=0)

{

    cerr<<"return from signal interruption";

    return 0;

}

else

    cerr<<"return from first time sigsetjmp is called";

pause();

}
```

KILL

- A process can send a signal to a related process via the kill API. This is a simple means of inter-process communication or control. The function prototype of the API is:

```
#include<signal.h>
```

```
int kill(pid_t pid, int signal_num);
```

Returns: 0 on success, -1 on failure.

The `signal_num` argument is the integer value of a signal to be sent to one or more processes designated by `pid`. The possible values of `pid` and its use by the `kill` API are:

<code>pid > 0</code>	The signal is sent to the process whose process ID is <code>pid</code> .
<code>pid == 0</code>	The signal is sent to all processes whose process group ID equals the process group ID of the sender and for which the sender has permission to send the signal.
<code>pid < 0</code>	The signal is sent to all processes whose process group ID equals the absolute value of <code>pid</code> and for which the sender has permission to send the signal.
<code>pid == 1</code>	The signal is sent to all processes on the system for which the sender has permission to send the signal.

The following program illustrates the implementation of the UNIX kill command using the kill API:

```
#include<iostream.h>
#include<stdio.h>
#include<unistd.h>
#include<string.h>
#include<signal.h>

int main(int argc,char** argv)
{
    int pid, sig = SIGTERM;
    if(argc==3)
    {
        if(sscanf(argv[1],"%d",&sig)!=1)
        {
            cerr<<"invalid number:" << argv[1] << endl;
            return -1;
        }
        argv++,argc--;
    }
    while(--argc>0)
    if(sscanf(*++argv, "%d", &pid)==1)
    {
        if(kill(pid,sig)==-1)
            perror("kill");
    }
    else
        cerr<<"invalid pid:" << argv[0] <<endl;
    return 0;
```

- The UNIX kill command invocation syntax is:

Kill [-<signal_num>] <pid>.....

where signal_num can be an integer number or the symbolic name of a signal. <pid> is process ID.

ALARM

The alarm API can be called by a process to request the kernel to send the SIGALRM signal after a certain number of real clock seconds.

The function prototype of the API is:

```
#include<signal.h>
```

```
Unsigned int alarm(unsigned int time_interval);
```

Returns: 0 or number of seconds until previously set alarm

The alarm API can be used to implement the sleep API:


```
#include<signal.h>
#include<stdio.h>
#include<unistd.h>
```

```
void wakeup( )
{ ; }
```

```
unsigned int sleep (unsigned int timer )
{
}
```

```
struct sigaction action;
action.sa_handler=wakeup;
action.sa_flags=0;
sigemptyset(&action.sa_mask);
if(sigaction(SIGALARM,&action,0)==-1)
{
perror("sigaction");
return -1;
}
```

```
(void) alarm (timer);
(void) pause( );
return 0;
```

INTERVAL TIMERS

- The interval timer can be used to schedule a process to do some tasks at a fixed time interval, to time the execution of some operations, or to limit the time allowed for the execution of some tasks.
- The following program illustrates how to set up a real-time clock interval timer using the alarm API:

```
#include<stdio.h>
#include<unistd.h>
#include<signal.h>
#define INTERVAL 5
void callme(int sig_no)
{
    alarm(INTERVAL);
    /*do scheduled tasks*/
}
int main()

{
    struct sigaction action;
    sigemptyset(&action.sa_mask);
    action.sa_handler=(void*)( ) callme; action.sa_flags=SA_RESTART;
    if(sigaction(SIGALARM,&action,0)==-1)
    {
    }
    perror("sigaction");
    return 1;
}
if(alarm(INTERVAL)==-1)
    perror("alarm");
else while(1)
{
    /*do normal operation*/
}
return 0;
```

- In addition to alarm API, UNIX also invented the setitimer API, which can be used to define up to three different types of timers in a process:
 - Real time clock timer
 - Timer based on the user time spent by a process
 - Timer based on the total user and system times spent by a process
- The getitimer API is also defined for users to query the timer values that are set by the setitimer API.

The setitimer and getitimer function prototypes are:

```
#include<sys/time.h>
```

```
int setitimer(int which, const struct itimerval * val, struct  
itimerval * old);
```

```
int getitimer(int which, struct itimerval * old);
```

The which arguments to the above APIs specify which timer to process. Its possible values and the corresponding timer types are:

ITIMER_REAL	decrements in real time and generates a SIGALRM signal when it expires
ITIMER_VIRTUAL	decrements in virtual time (time used by the process) and generates a SIGVTALRM signal when it expires.
ITIMER_PROF	decrements in virtual time and system time for the process and generates a SIGPROF signal when it expires.

- The struct itimerval datatype is defined as:

```
struct itimerval
```

```
{
```

```
};
```

```
struct timeval it_value; /*current value*/
```

```
struct timeval it_interval; /* time interval*/
```

- The setitimer and getitimer APIs return a zero value if they succeed or a -1 value if they fail.

Example program:

```
#include<stdio.h>
#include<unistd.h>
#include<signal.h>
#define INTERVAL 5
void callme(int sig_no)
{
    /*do scheduled tasks*/
}
int main()
{
    struct itimerval val;
    struct sigaction action;

    sigemptyset(&action.sa_mask);
    action.sa_handler=(void(*)()) callme;
    action.sa_flags=SA_RESTART;
    if(sigaction(SIGALARM,&action,0)==-1)
    {
        perror("sigaction");
        return 1;
    }
```



```
val.it_interval.tv_sec =INTERVAL;  
val.it_interval.tv_usec =0;  
val.it_value.tv_sec =INTERVAL;  
val.it_value.tv_usec =0;  
  
}  
  
if(setitimer(ITIMER_REAL, &val , 0)==-1)  
    perror(“alarm”);  
  
else while(1)  
{  
/*do normal operation*/  
}  
return 0;
```

POSIX.1b TIMERS

- POSIX.1b defines a set of APIs for interval timer manipulations. The POSIX.1b timers are more flexible and powerful than are the UNIX timers in the following ways:
 - Users may define multiple independent timers per system clock.
 - The timer resolution is in nanoseconds.
 - Users may specify the signal to be raised when a timer expires.
 - The time interval may be specified as either an absolute or a relative time.

The POSIX.1b APIs for timer manipulations are:

```
#include<signal.h>
#include<time.h>

int timer_create(clockid_t clock, struct sigevent* spec, timer_t* timer_hdrp);
int timer_settime(timer_t timer_hdr, int flag, struct itimerspec* val, struct itimerspec* old);
int timer_gettime(timer_t timer_hdr, struct itimerspec* old);
int timer_getoverrun(timer_t timer_hdr);
int timer_delete(timer_t timer_hdr);
```

DAEMON PROCESSES

INTRODUCTION

Daemons are processes that live for a long time. They are often started when the system is bootstrapped and terminate only when the system is shut down.

DAEMON CHARACTERISTICS

- The characteristics of daemons are:
 - ☐ Daemons run in background.
 - ☐ Daemons have super-user privilege.
 - ☐ Daemons don't have controlling terminal.
 - ☐ Daemons are session and group leaders.

CODING RULES

- ❑ **Call umask to set the file mode creation mask to 0.** The file mode creation mask that's inherited could be set to deny certain permissions. If the daemon process is going to create files, it may want to set specific permissions.
- ❑ **Call fork and have the parent exit.** This does several things. First, if the daemon was started as a simple shell command, having the parent terminate makes the shell think that the command is done. Second, the child inherits the process group ID of the parent but gets a new process ID, so we're guaranteed that the child is not a process group leader.
- ❑ **Call setsid to create a new session.** The process (a) becomes a session leader of a new session, (b) becomes the process group leader of a new process group, and (c) has no controlling terminal.

- **Change the current working directory to the root directory.**

The current working directory inherited from the parent could be on a mounted file system. Since daemons normally exist until the system is rebooted, if the daemon stays on a mounted file system, that file system cannot be unmounted.

- **Unneeded file descriptors should be closed.**

This prevents the daemon from holding open any descriptors that it may have inherited from its parent.

- **Some daemons open file descriptors 0, 1, and 2 to /dev/null so that any library routines that try to read from standard input or write to standard output or standard error will have no effect.**

Since the daemon is not associated with a terminal device, there is nowhere for output to be displayed; nor is there anywhere to receive input from an interactive user. Even if the daemon was started from an interactive session, the daemon runs in the background, and the login session can terminate without affecting the daemon.

If other users log in on the same terminal device, we wouldn't want output from the daemon showing up on the terminal, and the users wouldn't expect their input to be read by the daemon.

Example Program:

```
#include <unistd.h>
```

```
#include <sys/types.h>
```

```
#include <fcntl.h>
```

```
int daemon_initialise()
```

```
{
```

```
pid_t pid;
```

```
if (( pid = fork() ) < 0)
```

```
    return -1;
```

```
else if ( pid != 0)
```

```
    exit(0);    /* parent exits */
```

```
/* child continues */
```

```
setsid( );
```

```
chdir("/");
```

```
umask(0);
```

```
return 0;
```

```
}
```