

FILTERS USING REGULAR EXPRESSIONS – grep and sed

We often need to search a file for a pattern, either to see the lines containing (or not containing) it or to have it replaced with something else. This chapter discusses two important filters that are specially suited for these tasks – grep and sed. grep takes care of all search requirements we may have. sed goes further and can even manipulate the individual characters in a line. In fact sed can do several things, some of them quite well.

grep – searching for a pattern

It scans the file / input for a pattern and displays lines containing the pattern, the line numbers or filenames where the pattern occurs. It's a command from a special family in UNIX for handling search requirements.

```
grep options pattern filename(s)
```

```
grep "sales" emp.lst
```

will display lines containing sales from the file emp.lst. Patterns with and without quotes is possible. It's generally safe to quote the pattern. Quote is mandatory when pattern involves more than one word. It returns the prompt in case the pattern can't be located.

```
grep president emp.lst
```

When grep is used with multiple filenames, it displays the filenames along with the output.

```
grep "director" emp1.lst emp2.lst
```

Where it shows filename followed by the contents

grep options

grep is one of the most important UNIX commands, and we must know the options that POSIX requires grep to support. Linux supports all of these options.

-i	ignores case for matching
-v	doesn't display lines matching expression
-n	displays line numbers along with lines
-c	displays count of number of occurrences
-l	displays list of filenames only
-e exp	specifies expression with this option
-x	matches pattern with entire line
-f file	takes patterns from file, one per line
-E	treats pattern as an extended RE
-F	matches multiple fixed strings

```
grep -i 'agarwal' emp.lst
```

```
grep -v 'director' emp.lst > otherlist
```

wc -l otherlist will display 11 otherlist

```
grep -n 'marketing' emp.lst
```

```
grep -c 'director' emp.lst
```

```
grep -c 'director' emp*.lst
```

will print filenames prefixed to the line count

```
grep -l 'manager' *.lst
```

will display filenames *only*

```
grep -e 'Agarwal' -e 'aggarwal' -e 'agrawal' emp.lst
```

will print matching multiple patterns

```
grep -f pattern.lst emp.lst
```

all the above three patterns are stored in a separate file *pattern.lst*

Basic Regular Expressions (BRE) – An Introduction

It is tedious to specify each pattern separately with the -e option. grep uses an expression of a different type to match a group of similar patterns. If an expression uses meta characters, it is termed a regular expression. Some of the characters used by regular expression are also meaningful to the shell.

BRE character subset

The basic regular expression character subset uses an elaborate meta character set, overshadowing the shell's wild-cards, and can perform amazing matches.

*	Zero or more occurrences
g*	nothing or g, gg, ggg, etc.
.	A single character
.*	nothing or any number of characters
[pqr]	a single character p, q or r
[c1-c2]	a single character within the ASCII range represented by c1 and c2

The character class

grep supports basic regular expressions (BRE) by default and extended regular expressions (ERE) with the `-E` option. A regular expression allows a group of characters enclosed within a pair of `[]`, in which the match is performed for a single character in the group.

```
grep "[aA]g[ar][ar]wal" emp.lst
```

A single pattern has matched two similar strings. The pattern `[a-zA-Z0-9]` matches a single alphanumeric character. When we use range, make sure that the character on the left of the hyphen has a lower ASCII value than the one on the right. Negating a class (`^`) (caret) can be used to negate the character class. When the character class begins with this character, all characters other than the ones grouped in the class are matched.

The *

The asterisk refers to the immediately preceding character. `*` indicates zero or more occurrences of the previous character.

`g*` nothing or `g`, `gg`, `ggg`, etc.

```
grep "[aA]gg*[ar][ar]wal" emp.lst
```

Notice that we don't require to use `-e` option three times to get the same output!!!!

The dot

A dot matches a single character. The shell uses `?` Character to indicate that.

`.*` signifies any number of characters or none

```
grep "j.*saxena" emp.lst
```

Specifying Pattern Locations (^ and \$)

Most of the regular expression characters are used for matching patterns, but there are two that can match a pattern at the beginning or end of a line. Anchoring a pattern is often necessary when it can occur in more than one place in a line, and we are interested in its occurrence only at a particular location.

<code>^</code>	for matching at the beginning of a line
<code>\$</code>	for matching at the end of a line

```
grep "^2" emp.lst
```

Selects lines where emp_id starting with 2

```
grep "7...$" emp.lst
```

Selects lines where emp_salary ranges between 7000 to 7999

```
grep "^[^2]" emp.lst
```

Selects lines where emp_id doesn't start with 2

When meta characters lose their meaning

It is possible that some of these special characters actually exist as part of the text. Sometimes, we need to escape these characters. For example, when looking for a pattern `g*`, we have to use `\`
To look for `[`, we use `\[`
To look for `.*`, we use `\.*`

Extended Regular Expression (ERE) and grep

If current version of grep doesn't support ERE, then use egrep but without the `-E` option. `-E` option treats pattern as an ERE.

`+` matches one or more occurrences of the previous character

`?` Matches zero or one occurrence of the previous character

`b+` matches `b`, `bb`, `bbb`, etc.

`b?` matches either a single instance of `b` or nothing

These characters restrict the scope of match as compared to the `*`

```
grep -E "[aA]gg?arwal" emp.lst
```

```
# ?include <stdio.h>
```

The ERE set

<code>ch+</code>	matches one or more occurrences of character <code>ch</code>
<code>ch?</code>	Matches zero or one occurrence of character <code>ch</code>
<code>exp1 exp2</code>	matches <code>exp1</code> or <code>exp2</code>
<code>(x1 x2)x3</code>	matches <code>x1x3</code> or <code>x2x3</code>

Matching multiple patterns (`|`, `(` and `)`)

```
grep -E 'sengupta|dasgupta' emp.lst
```

We can locate both without using `-e` option twice, or

```
grep -E '(sen|das)gupta' emp.lst
```

sed – The Stream Editor

sed is a multipurpose tool which combines the work of several filters. sed uses **instructions** to act on text. An instruction combines an **address** for selecting lines, with an **action** to be taken on them.

```
sed options 'address action' file(s)
```

sed supports only the BRE set. *Address* specifies either one line number to select a single line or a set of two lines, to select a group of contiguous lines. *action* specifies print, insert, delete, substitute the text.

sed processes several instructions in a sequential manner. Each instruction operates on the output of the previous instruction. In this context, two options are relevant, and probably they are the only ones we will use with sed – the `-e` option that lets us use multiple instructions, and the `-f` option to take instructions from a file. Both options are used by grep in identical manner.

Line Addressing

```
sed '3q' emp.lst
```

Just similar to *head -n 3 emp.lst*. Selects first three lines and quits

```
sed -n '1,2p' emp.lst
```

`p` prints selected lines as well as all lines. To suppress this behavior, we use `-n` whenever we use `p` command

```
sed -n '$p' emp.lst
```

Selects last line of the file

```
sed -n '9,11p' emp.lst
```

Selecting lines from anywhere of the file, between lines from 9 to 11

```
sed -n '1,2p  
7,9p  
$p' emp.lst
```

Selecting multiple groups of lines

```
sed -n '3,$!p' emp.lst
```

Negating the action, just same as 1,2p

Using Multiple Instructions (-e and -f)

There is adequate scope of using the -e and -f options whenever sed is used with multiple instructions.

```
sed -n -e '1,2p' -e '7,9p' -e '$p' emp.lst
```

Let us consider,

```
cat instr.fil  
1,2p  
7,9p  
$p
```

-f option to direct the sed to take its instructions from the file

```
sed -n -f instr.fil emp.lst
```

We can combine and use -e and -f options as many times as we want

```
sed -n -f instr.fil1 -f instr.fil2 emp.lst
```

```
sed -n -e '/saxena/p' -f instr.fil1 -f instr.fil2 emp.lst
```

Context Addressing

We can specify one or more patterns to locate lines

```
sed -n '/director/p' emp.lst
```

We can also specify a comma-separated pair of context addresses to select a group of lines.

```
sed -n '/dasgupta/,/saxena/p' emp.lst
```

Line and context addresses can also be mixed

```
sed -n '1,/dasgupta/p' emp.lst
```

Using regular expressions

Context addresses also uses regular expressions.

```
Sed -n '/[aA]gg*[ar][ar]wal/p' emp.lst
```

Selects all agarwals.

```
Sed -n '/sa[kx]s*ena/p  
/gupta/p' emp.lst
```

Selects saxenas and gupta.

We can also use ^ and \$, as part of the regular expression syntax.

```
sed -n '/50.....$/p' emp.lst
```

Selects all people born in the year 1950.

Writing Selected Lines to a File (w)

We can use w command to write the selected lines to a separate file.

```
sed -n '/director/w dlist' emp.lst
```

Saves the lines of directors in *dlist* file

```
sed -n '/director/w dlist  
/manager/w mlist  
/executive/w elist' emp.lst
```

Splits the file among three files

```
sed -n '1,500w foo1  
501,$w foo2' foo.main
```

Line addressing also is possible. Saves first 500 lines in foo1 and the rest in foo2

Text Editing

sed supports inserting (i), appending (a), changing (c) and deleting (d) commands for the text.

```
$ sed '1\  
> #include <stdio.h>\  
> #include <unistd.h>  
> 'foo.c > $$
```

Will add two include lines in the beginning of foo.c file. Sed identifies the line without the \ as the last line of input. Redirected to \$\$ temporary file. This technique has to be followed when using the a and c commands also. To insert a blank line *after* each line of the file is printed (*double spacing text*), we have,

```
sed 'a\  
  
' emp.lst
```

Deleting lines (d)

```
sed '/director/d' emp.lst > olist      or  
  
sed -n '/director/!p' emp.lst > olist
```

Selects all lines except those containing *director*, and saves them in *olist*

Note that -n option not to be used with d

Substitution (s)

Substitution is the most important feature of sed, and this is one job that sed does exceedingly well.

```
[address]s/expression1/expression2/flags
```

Just similar to the syntax of substitution in vi editor, we use it in sed also.

```
sed 's/|/:/' emp.lst | head -n 2
```

```
2233:a.k.shukla |gm |sales |12/12/52|6000
```

```
9876:jai sharma |director|production|12/03/50|7000
```

Only the first instance of | in a line has been replaced. We need to use the g (global) flag to replace all the pipes.

```
sed 's/|/:/g' emp.lst | head -n 2
```

We can limit the vertical boundaries too by specifying an address (for first three lines only).

```
sed '1,3s/|/:/g' emp.lst
```

Replace the word director with member in the first five lines of emp.lst

```
sed '1,5s/director/member/' emp.lst
```


sed also uses regular expressions for patterns to be substituted. To replace all occurrence of agarwal, aggarwal and agrawal with simply Agarwal, we have,

```
sed 's/[Aa]gg*[ar][ar]wal/Agarwal/g' emp.lst
```

We can also use ^ and \$ with the same meaning. To add 2 prefix to all emp-ids,

```
sed 's/^/2/' emp.lst | head -n 1
```

```
22233 | a.k.shukla | gm | sales | 12/12/52 | 6000
```

To add .00 suffix to all salary,

```
sed 's/$/.00/' emp.lst | head -n 1
```

```
2233 | a.k.shukla | gm | sales | 12/12/52 | 6000.00
```

Performing multiple substitutions

```
sed 's/<I>/<EM>/g  
s/<B>/<STRONG>/g  
s/<U>/<EM>/g' form.html
```

An instruction processes the output of the previous instruction, as sed is a stream editor and works on data stream

```
sed 's/<I>/<EM>/g  
s/<EM>/<STRONG>/g' form.html
```

When a 'g' is used at the end of a substitution instruction, the change is performed globally along the line. Without it, only the left most occurrence is replaced. When there are a group of instructions to execute, you should place these instructions in a file instead and use sed with the -f option.

Compressing multiple spaces

```
sed 's/*|/|/g' emp.lst | tee empn.lst | head -n 3  
  
2233|a.k.shukla|g.m|sales|12/12/52|6000  
9876|jai sharma|director|production|12/03/50|7000  
5678|sumit chakrobarty|dgm|mrking|19/04/43|6000
```

The remembered patterns

Consider the below three lines which does the same job

```
sed 's/director/member/' emp.lst
```

```
sed '/director/s//member/' emp.lst
```

```
sed '/director/s/director/member/' emp.lst
```

The // representing an empty regular expression is interpreted to mean that the search and substituted patterns are the same

```
sed 's//g' emp.lst
```

 removes every | from file

Basic Regular Expressions (BRE) – Revisited

Three more additional types of expressions are:

- The repeated patterns - &

- The interval regular expression (IRE) – { }

- The tagged regular expression (TRE) – ()

The repeated patterns - &

To make the entire source pattern appear at the destination also

```
sed 's/director/executive director/' emp.lst
```

```
sed 's/director/executive &/' emp.lst
```

```
sed '/director/s//executive &/' emp.lst
```

Replaces director with executive director where & is a repeated pattern

The interval RE - { }

sed and grep uses IRE that uses an integer to specify the number of characters preceding a pattern. The IRE uses an escaped pair of curly braces and takes three forms:

- ch\{m\} – the ch can occur m times

- ch\{m,n\} – ch can occur between m and n times

- ch\{m,\} – ch can occur at least m times

The value of m and n can't exceed 255. Let teledir.txt maintains landline and mobile phone numbers. To select only mobile numbers, use IRE to indicate that a numerical can occur 10 times.

```
grep '[0-9]\{10\}' teledir.txt
```

Line length between 101 and 150

```
grep '^.\{101,150\}$' foo
```

Line length at least 101

```
sed -n '/.{101,}/p' foo
```

The Tagged Regular Expression (TRE)

You have to identify the segments of a line that you wish to extract and enclose each segment with a matched pair of escaped parenthesis. If we need to extract a number, `\([0-9]*\)`. If we need to extract non alphabetic characters,

```
\([^a-zA-Z]*\)
```

Every grouped pattern automatically acquires the numeric label `n`, where `n` signifies the `n`th group from the left.

```
sed 's/\ (a-z)*\ * \ ([a-z]*\) /\2, \1/' teledir.txt
```

To get surname first followed by a `,` and then the name and rest of the line. `sed` does not use compulsorily a `/` to delimit patterns for substitution. We can use only any character provided it doesn't occur in the entire command line. Choosing a different delimiter has allowed us to get away without escaping the `/` which actually occurs in the pattern.

-
- Source: Sumitabha Das, “UNIX – Concepts and Applications”, 4th edition, Tata McGraw Hill, 2006